



Database Technology

Indexing with R-trees

Motivation

In many real-life applications objects are represented as multidimensional points:

Spatial databases

(e.g., points in 2D or 3D space)


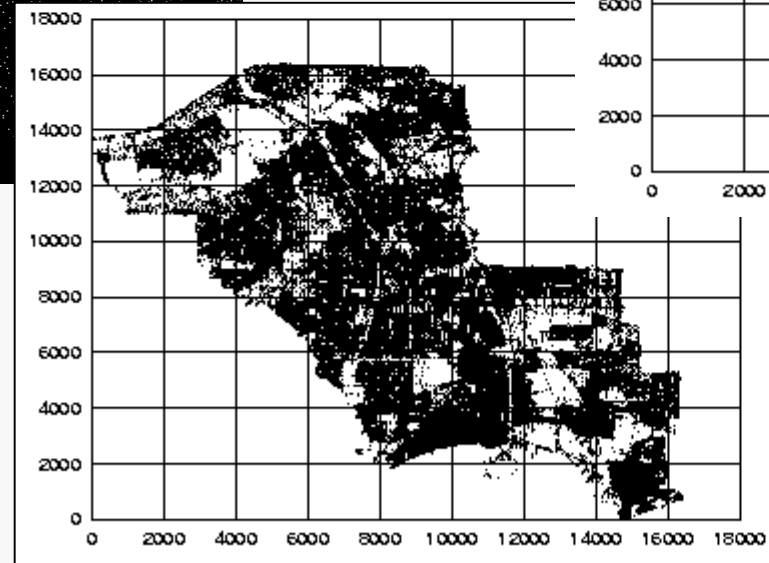

Multidimensional databases

(each record is considered a point in n -D space)

Multimedia databases

(e.g., feature vectors are extracted from images, audio files)

Examples of 2D Data Sets



Requirements

- Indexing scheme are needed to speed up query processing.
- We need disk-based techniques, since we do not want to be constrained by the memory capacity.
- The methods should handle insertions/deletions of objects (i.e., they should work in a dynamic environment).



The R-tree

A. Guttman:

“R-tree: A Dynamic Index Structure for Spatial Searching”,
ACM SIGMOD Conference, 1984

R-tree Index Structure

- An R-tree is a **height-balanced** tree similar to a **B-tree**.
- Index records in its leaf nodes containing pointers to data objects.
- Nodes correspond to disk pages if the index is disk-resident.
- The index is completely dynamic.
- **Leaf** node structure ($r, objectID$)
 - r : minimum bounding rectangle of object
 - $objectID$: the identifier of the corresponding object
- **Nonleaf** node structure ($R, childPTR$)
 - R : covers all rectangles in the lower node
 - $childPTR$: the address of a lower node in the R-tree

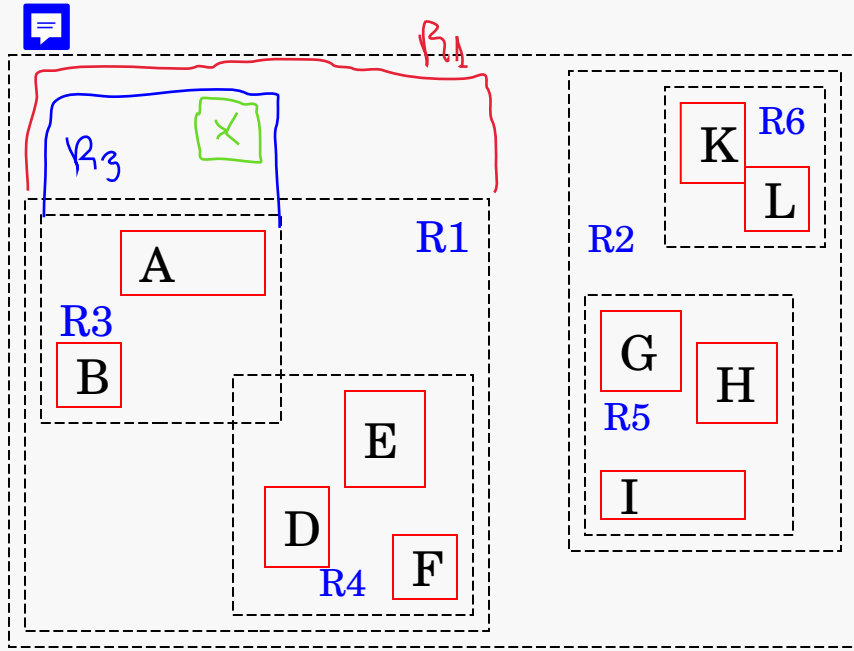
Properties of the R-tree

M : maximum number of entries that will fit in one node

m : minimum number of entries in a node, $m \leq M/2$

- Every leaf node contains between m and M index records unless it is the root.
- For each index record $(r, objectID)$ in a leaf node, r is the smallest rectangle (Minimum Bounding Rectangle (MBR)) that spatially contains the data object.
- Every non-leaf node has between m and M children, unless it is the root.
- For each entry $(R, childPTR)$ in a non-leaf node, R is the smallest rectangle that spatially contains the rectangles in the child node.
- The root node has at least 2 children unless it is a leaf.
- All leaves appear at the same level.


R-Tree Example



M : maximum number of entries

m : minimum number of entries ($\leq M/2$)

- (1) Every leaf node contains between m and M index records unless it is the root.
- (2) Each leaf node has the smallest rectangle that spatially contains the n -dimensional data objects.
- (3) Every non-leaf node has between m and M children unless it is the root.
- (4) Each non-leaf node has the smallest rectangle that spatially contains the rectangles in the child node.
- (5) The root node has at least two children unless it is a leaf.
- (6) All leaves appear on the same level.



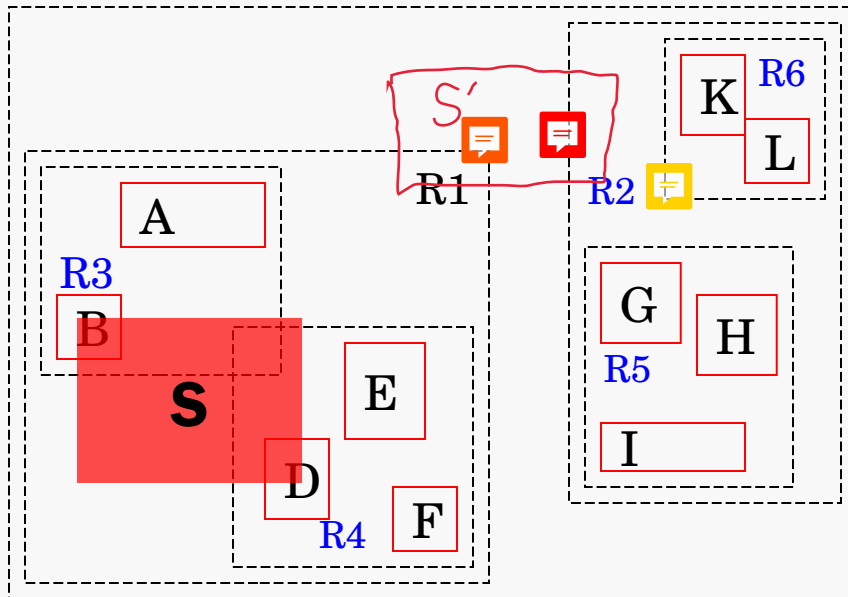
<MBR, Pointer to a child node>

<MBR, Pointer or ID>


R-tree Search Algorithm

- Given an R-tree whose root is T , find all index records whose rectangles overlap a search rectangle S .
- Algorithm Search
 - [Search subtrees]
 - If T is not a leaf, check each entry E to determine whether $E.R$ overlaps S .
 - For all overlapping entries, invoke Search on the tree whose root is pointed to by $E.childPTR$.
 - [Search leaf node]
 - If T is a leaf, check all entries E to determine whether $E.r$ overlaps S . If so, E is a qualifying record.


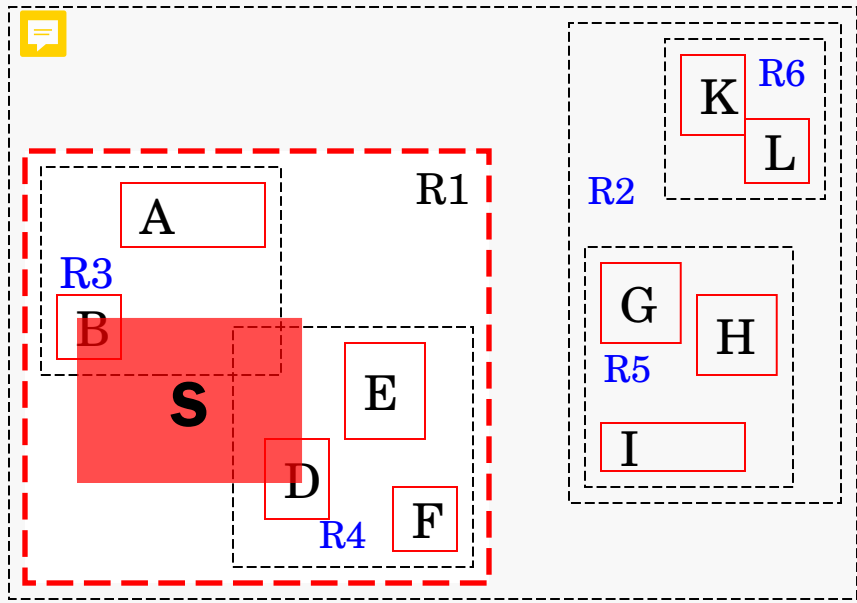
R-Tree Search Example (1/7)




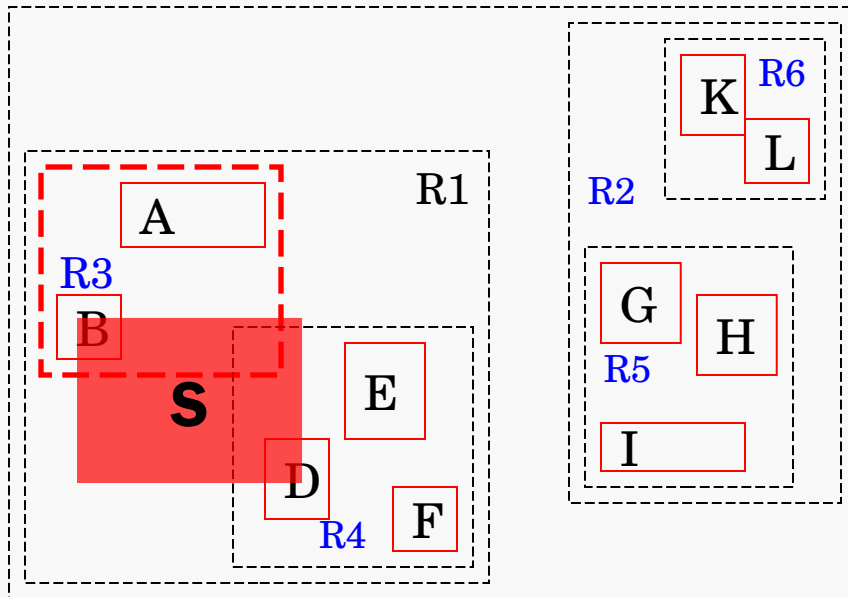
Find all objects whose rectangles are overlapped with a search rectangle S





R-Tree Search Example (2/7)




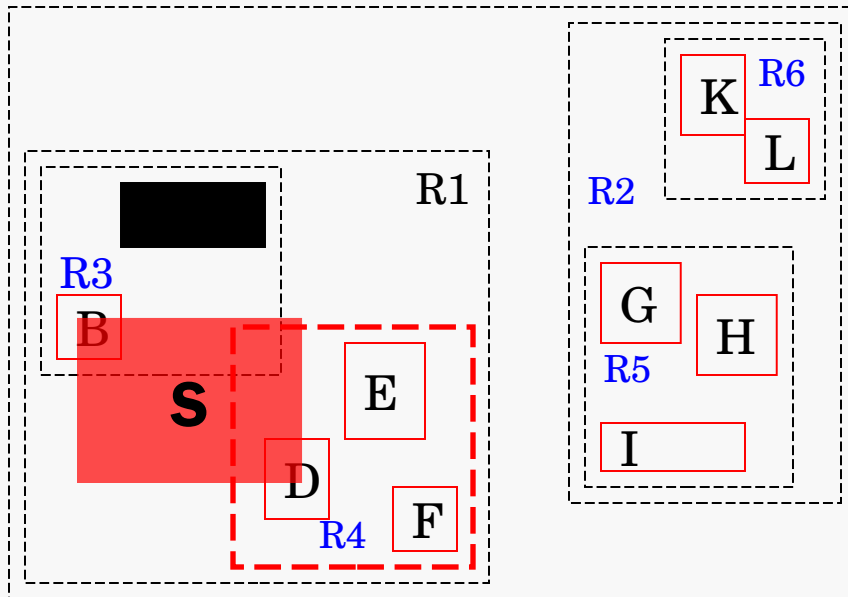
R-Tree Search Example (3/7)





R-Tree Search Example (4/7)





R-Tree Search Example (5/7)



R-Tree Search Example (6/7)



R-Tree Search Example (7/7)



B and D
→ overlapped objects with S



R-tree Insertion



➤ Algorithm **Insert**

// Insert a new index entry E into an R-tree.

(1)- [Find position for new record] 

- Invoke **ChooseLeaf** to select a leaf node **L** in which to place E.

(2)- [Add record to leaf node]

- If L has room for another entry, install E.
- Otherwise invoke **SplitNode** to obtain **L** and **LL** containing E and all the old entries of L.

(3)- [Propagate changes upward]

- Invoke **AdjustTree** on **L**, also passing **LL** if a split was performed

- [Grow tree taller]

- If node split propagation caused the root to split, create a **new root** whose children are the 2 resulting nodes.

Algorithm ChooseLeaf (1)



// Select a leaf node in which to place a new index entry E.

- [Initialize]
 - Set N to be the root node.
- [Leaf check]*
 - If N is not a leaf, return N.
- [Choose subtree]
 - If N is not a leaf, let F be the entry in N whose rectangle F.I needs least enlargement to include E.I.
Resolve ties by choosing the entry with the rectangle of smallest area.
- [Descend until a leaf is reached]
 - Set N to be the child node pointed to by F.p.
 - Repeat from *.

Algorithm AdjustTree {3}

- // Ascend from a leaf node **L** to the root, adjusting covering rectangles and propagating node splits as necessary.
- [Initialize]
 - Set $N=L$. If L was split previously, set NN to be the resulting second node.
 - [Check if done]*
 - If N is the root, stop.
 - [Adjust covering rectangle in parent entry]
 - Let P be the parent node of N , and let E_n be N 's entry in P .
 - Adjust $E_n.I$ so that it tightly encloses all entry rectangles in N .
 - [Propagate node split upward]
 - If N has a partner NN resulting from an earlier split, create a new entry E_{NN} with $E_{NN}.p$ pointing to NN and $E_{NN}.I$ enclosing all rectangles in NN .
 - Add E_{NN} to P if there is room. Otherwise, invoke **SplitNode** to produce P and PP containing E_{NN} and all P 's old entries.
 - [Move up to next level]
 - Set $N = P$ and set $NN = PP$ if a split occurred. Repeat from *.

R-tree Deletion

➤ Algorithm **Delete**

// Remove index record E from an R-tree.

- [Find node containing record]
 - Invoke **FindLeaf** to locate the leaf node L containing E.
 - Stop if the record was not found.
- [Delete record]
 - Remove E from L.
- [Propagate changes]
 - Invoke **CondenseTree**, passing L.
- [Shorten tree]
 - If the root node has only one child after the tree has been adjusted, make the child the new root.

➤ Algorithm **FindLeaf**

// Find the leaf node containing the entry E in an R-tree with root T.


- [Search subtrees]
 - If T is not a leaf, check each entry F in T to determine if **F.I** overlaps **E.I**. For each such entry invoke **FindLeaf** on the tree whose root is pointed to by F.p until E is found or all entries have been checked.
- [Search leaf node for record]
 - If T is a leaf, check each entry to see if it matches E. If E is found return T.

Algorithm CondenseTree


- // Given a leaf node L from which an entry has been deleted, eliminate the node if it
- // has too few entries and relocate its entries.
- // Propagate node elimination upward as necessary.
- // Adjust all covering rectangles on the path to the root.
- [Initialize]
 - Set $N=L$. Set Q , the set of eliminated nodes, to be empty.
- [Find parent entry]*
 - If N is the root, go to +.
 - Otherwise, let P be the parent of N , and let E_n be N 's entry in P .
- [Eliminate under-full node]
 - If N has fewer than m entries, delete E_n from P and add N to set Q .
- [Adjust covering rectangle]
 - If N has not been eliminated, adjust $E_n.I$ to tightly contain all entries in N .
- [Move up one level in tree]
 - Set $N=P$ and repeat from *.
- [Re-insert orphaned entries]+
 - Re-insert all entries of nodes in set Q .
Entries from eliminated leaf nodes are re-inserted in tree leaves as described in *Insert*, but entries from higher-level nodes must be placed higher in the tree, so that leaves of their dependent subtrees will be on the same level as leaves of the main tree.

Node Splitting

- The **total area** of the 2 covering rectangles after a split should be **minimized**.
 - ⇒ The same criterion was used in **ChooseLeaf** to decide a new index entry: at each level in the tree, the subtree chosen was the one whose covering rectangle would have to be **enlarged least**.



bad split



good split

Node Split Algorithms (2)

➤ Exhaustive Algorithm

- To generate all possible groupings and choose the best.
⇒ The number of possible splits is very large.

➤ Quadratic-Cost Algorithm

- Attempts to find a **small-area split**, but is not guaranteed to find one with the smallest area possible.
- Quadratic in M (node capacity) and linear in dimensionality
- Picks two of the $M+1$ entries to be the first elements of the 2 new groups by choosing the pair that would waste the most area if both were put in the same group, i.e., the area of a rectangle covering both entries would be greatest.
- The remaining entries are then assigned to groups one at a time.
- At each step the area expansion required to add each remaining entry to each group is calculated, and the entry assigned is the one showing the greatest difference between the 2 groups.

Algorithm Quadratic Split

// Divide a set of $M+1$ index entries into 2 groups.

- [Pick first entry for each group]
 - Apply algorithm **PickSeeds** to choose 2 entries to be the first elements of the groups.
 - Assign each to a group.
- [Check if done]*
 - If all entries have been assigned, stop.
 - If one group has so few entries that all the rest must be assigned to it in order for it to have the minimum number **m**, assign them and stop.
- [Select entry to assign]
 - Invoke algorithm **PickNext** to choose the next entry to assign.
 - Add it to the group whose covering rectangle will have to be **enlarged least** to accommodate it.
 - Resolve ties by adding the entry to the group with smaller entry, then to the one with fewer entries, then to either.
 - Repeat from *.

Algorithms PickSeeds & PickNext

➤ Algorithm PickSeeds

// Select 2 entries to be the first elements of the groups.


- [Calculate inefficiency of grouping entries together]
 - For each pair of entries E_1 and E_2 , compose a rectangle J including $E_1.I$ and $E_2.I$.
 - Calculate $d = \text{area}(J) - \text{area}(E_1.I) - \text{area}(E_2.I)$.
- [Choose the most wasteful pair.]
 - Choose the pair with the largest d .

➤ Algorithm PickNext

// Select one remaining entry for classification in a group.

- [Determine cost of putting each entry in each group]
 - For each entry E not yet in a group,
 - Calculate d_1 = the area increase required in the covering rectangle of Group 1 to include $E.I$.
 - Calculate d_2 similarly for Group 2.
- [Find entry with greatest preference for one group]
 - Choose any entry with the maximum difference between d_1 & d_2 .

A Linear-Cost Algorithm

- Linear in M and in dimensionality
- **Linear Split** is identical to Quadratic Split but uses a different **PickSeeds**. **PickNext** simply chooses any of the remaining entries.
- Algorithm **LinearPickSeeds** 
 - // Select 2 entries to be the first elements of the groups.
 - [Find extreme rectangles along all dimensions]
 - Along each dimension, find the entry whose rectangle has the highest low side, and the one with the lowest high side.
 - Record the separation.
 - [Adjust for shape of the rectangle cluster]
 - Normalize the separations by dividing by the width of the entire set along the corresponding dimension.
 - [Select the most extreme pair]
 - Choose the pair with the greatest normalized separation along any dimension.

Performance (Insert/Delete/Search)




Figure 4.7
CPU cost of inserts and deletes
vs. amount of data.




Figure 4.8
Search performance vs. amount of data:
Pages touched

Performance (Search/Space)




Figure 4.9
Search performance vs. amount of data:
CPU cost




Figure 4.10
Space required for R-tree
vs. amount of data.

Conclusions

- The R-tree structure has been shown to be useful for indexing spatial data objects of non-zero size.
- The **linear node-split** algorithm proved to be as good as more expensive techniques.
 - ⇒ It was fast, and the slightly worse quality of the splits did not affect search performance noticeably.